

# GEOG 4/590: Geospatial Data Science

## Lecture 2: Vector data analysis



Email: [jryan4@uoregon.edu](mailto:jryan4@uoregon.edu)

Office: 163A Condon Hall

Office hours: Monday 15:00-16:00 and Tuesday 14:00-15:00

# Vector data analysis


The vector data model represents space as a series of discrete entities such as borders, buildings, streets, and roads. There are three different types of vector data: points, lines and polygons. Online mapping applications, such as **Google Maps** and **OpenStreetMap**, use this format to display data.

# Vector data analysis

The vector data model represents space as a series of discrete entities such as borders, buildings, streets, and roads. There are three different types of vector data: points, lines and polygons. Online mapping applications, such as **Google Maps** and **OpenStreetMap**, use this format to display data.

The Python library **GeoPandas** provides some great tools for working with vector data. As the name suggests, **GeoPandas** extends the popular data science library **Pandas** by adding support for geospatial data. The core data structure in **GeoPandas** is the **GeoDataFrame**. The key difference between the two is that a **GeoDataFrame** can store geometry data and perform spatial operations.

# Vector data analysis



The **geometry** column can contain any geometry type (e.g. points, lines, polygons) or even a mixture.

# Reading files

Assuming we have a file containing both data and geometry (e.g. GeoPackage, GeoJSON, Shapefile), we can read it using `read_file`, which automatically detects the filetype and creates a `GeoDataFrame`. In the this demo, we will be working with three shapefiles containing 1) cities and towns (as points), 2) urban growth boundaries (as polygons), and 3) counties (as polygons) in Oregon.

# Reading files

Assuming we have a file containing both data and geometry (e.g. GeoPackage, GeoJSON, Shapefile), we can read it using `read_file`, which automatically detects the filetype and creates a `GeoDataFrame`. In the this demo, we will be working with three shapefiles containing 1) cities and towns (as points), 2) urban growth boundaries (as polygons), and 3) counties (as polygons) in Oregon.

```
import geopandas as gpd

cities = gpd.read_file('data/oregon_cities.shp')
cities.head()
```




# Reading files

Assuming we have a file containing both data and geometry (e.g. GeoPackage, GeoJSON, Shapefile), we can read it using `read_file`, which automatically detects the filetype and creates a `GeoDataFrame`. In the this demo, we will be working with three shapefiles containing 1) cities and towns (as points), 2) urban growth boundaries (as polygons), and 3) counties (as polygons) in Oregon.

```
import geopandas as gpd

cities = gpd.read_file('data/oregon_cities.shp')
cities.head()
```

	name	lat	lon	geometry
0	Adair Village city	44.67	-123.22	POINT (-123.22000 44.67000)
1	Adams	45.77	-118.56	POINT (-118.56000 45.77000)
2	Adrian	43.74	-117.07	POINT (-117.07000 43.74000)
3	Albany	44.63	-123.10	POINT (-123.10000 44.63000)
4	Aloha	45.49	-122.87	POINT (-122.87000 45.49000)



# DataFrame properties

We can analyze our **GeoDataFrame** using standard **Pandas** functions.

```
# Data types of each column  
cities.dtypes
```

```
name           object  
lat            float64  
lon            float64  
geometry       geometry  
dtype: object
```

	name	lat	lon	geometry
0	Adair Village city	44.67	-123.22	POINT (-123.22000 44.67000)
1	Adams	45.77	-118.56	POINT (-118.56000 45.77000)
2	Adrian	43.74	-117.07	POINT (-117.07000 43.74000)
3	Albany	44.63	-123.10	POINT (-123.10000 44.63000)
4	Aloha	45.49	-122.87	POINT (-122.87000 45.49000)



# DataFrame properties

We can analyze our **GeoDataFrame** using standard **Pandas** functions.

```
# Number of rows and columns  
cities.shape
```

```
(377, 4)
```

	name	lat	lon	geometry
0	Adair Village city	44.67	-123.22	POINT (-123.22000 44.67000)
1	Adams	45.77	-118.56	POINT (-118.56000 45.77000)
2	Adrian	43.74	-117.07	POINT (-117.07000 43.74000)
3	Albany	44.63	-123.10	POINT (-123.10000 44.63000)
4	Aloha	45.49	-122.87	POINT (-122.87000 45.49000)

# DataFrame properties

We can analyze our **GeoDataFrame** using standard **Pandas** functions.

```
# Name of columns  
cities.columns
```

```
Index(['name', 'lat', 'lon', 'geometry'], dtype='object')
```

	name	lat	lon	geometry
0	Adair Village city	44.67	-123.22	POINT (-123.22000 44.67000)
1	Adams	45.77	-118.56	POINT (-118.56000 45.77000)
2	Adrian	43.74	-117.07	POINT (-117.07000 43.74000)
3	Albany	44.63	-123.10	POINT (-123.10000 44.63000)
4	Aloha	45.49	-122.87	POINT (-122.87000 45.49000)

# Indexing

We can select specific columns based on the column values. The basic syntax is `dataframe[value]`, where `value` can be a single column name, or a list of column names.

```
# List the city names
cities['name']
```

```
0      Adair Village city
1              Adams
2              Adrian
3              Albany
4              Aloha
...
372      Wood Village
373      Woodburn
374      Yachats
375      Yamhill
376      Yoncalla
Name: name, Length: 377, dtype: object
```

	name	lat	lon	geometry
0	Adair Village city	44.67	-123.22	POINT (-123.22000 44.67000)
1	Adams	45.77	-118.56	POINT (-118.56000 45.77000)
2	Adrian	43.74	-117.07	POINT (-117.07000 43.74000)
3	Albany	44.63	-123.10	POINT (-123.10000 44.63000)
4	Aloha	45.49	-122.87	POINT (-122.87000 45.49000)

# Indexing

We can select specific columns based on the column values. The basic syntax is `dataframe[value]`, where `value` can be a single column name, or a list of column names.

```
# List the latitudes and longitudes
cities[['lat','lon']]
```

	lat	lon
0	44.67	-123.22
1	45.77	-118.56
2	43.74	-117.07
3	44.63	-123.10
4	45.49	-122.87
...	...	...
372	45.54	-122.42
373	45.15	-122.86
374	44.31	-124.10
375	45.34	-123.19
376	43.60	-123.29

377 rows x 2 columns

# Indexing

We can select specific rows using the `.iloc` method.

```
# Second row  
cities.iloc[1]
```

```
name          Adams  
lat           45.77  
lon          -118.56  
geometry      POINT (-118.56 45.77)  
Name: 1, dtype: object
```

# Indexing

We can select specific rows using the `.iloc` method.

```
# Sixth to tenth rows  
cities.iloc[5:10]
```

	name	lat	lon	geometry
5	Alpine	44.33	-123.36	POINT (-123.36000 44.33000)
6	Alsea	44.38	-123.60	POINT (-123.60000 44.38000)
7	Altamont	42.20	-121.72	POINT (-121.72000 42.20000)
8	Amity	45.12	-123.20	POINT (-123.20000 45.12000)
9	Annex	44.23	-116.99	POINT (-116.99000 44.23000)

# Masking

We can sample of our **DataFrame** based on specific values by producing a **Boolean mask** (i.e. a list of values equal to **True** or **False**). To find cities that are East of -117.5 degrees longitude, we could write:

```
mask = cities['lon'] > -117.5  
cities[mask]
```

# Masking

We can sample of our **DataFrame** based on specific values by producing a **Boolean mask** (i.e. a list of values equal to **True** or **False**). To find cities that are East of -117.5 degrees longitude, we could write:

```
mask = cities['lon'] > -117.5  
cities[mask]
```

	name	lat	lon	geometry
2	Adrian	43.74	-117.07	POINT (-117.07000 43.74000)
9	Annex	44.23	-116.99	POINT (-116.99000 44.23000)
97	Enterprise	45.43	-117.28	POINT (-117.28000 45.43000)
134	Halfway	44.88	-117.11	POINT (-117.11000 44.88000)
150	Huntington	44.35	-117.27	POINT (-117.27000 44.35000)
164	Jordan Valley	42.98	-117.06	POINT (-117.06000 42.98000)
165	Joseph	45.35	-117.23	POINT (-117.23000 45.35000)
190	Lostine	45.49	-117.43	POINT (-117.43000 45.49000)



# Masking

It's more concise to just add the Boolean mask between square brackets. Here we find cities that contain a **z** in their name.

```
cities[cities['name'].str.contains('z')]
```

	name	lat	lon	geometry
<b>34</b>	Bonanza	42.20	-121.41	POINT (-121.41000 42.20000)
<b>168</b>	Keizer	45.00	-123.02	POINT (-123.02000 45.00000)
<b>195</b>	Manzanita	45.72	-123.94	POINT (-123.94000 45.72000)
<b>206</b>	Metzger	45.45	-122.76	POINT (-122.76000 45.45000)
<b>302</b>	Siletz	44.72	-123.92	POINT (-123.92000 44.72000)

# Masking

Or use string matching to find a specific city.

```
cities[cities['name'] == 'Eugene']
```

	name	lat	lon	geometry
100	Eugene	44.06	-123.12	POINT (-123.12000 44.06000)

# Descriptive statistics

**Pandas** provides basic functions to calculate descriptive statistics.

```
# Minimum latitude value  
cities['lat'].min()
```

```
42.0
```

# Descriptive statistics

**Pandas** provides basic functions to calculate descriptive statistics.

```
# Minimum latitude value  
cities['lat'].min()
```

```
42.0
```

```
# Mean longitude value  
cities['lon'].mean()
```

```
-122.02392572944296
```

# Descriptive statistics

**Pandas** provides basic functions to calculate descriptive statistics.

Sometimes we want to know which row contains the specific value which we can do using **idxmax/idxmin**.

```
cities['lat'].idxmin()
```

```
232
```

# Descriptive statistics

**Pandas** provides basic functions to calculate descriptive statistics.

Sometimes we want to know which row contains the specific value which we can do using **idxmax/idxmin**.

```
cities['lat'].idxmin()
```

```
232
```

```
cities.iloc[232]
```

```
name          New Pine Creek
lat           42.0
lon          -120.3
geometry      POINT (-120.3 42)
Name: 232, dtype: object
```

# Sorting

We can sort `DataFrames` using the `sort_values` function. This function takes two arguments, `by` and `ascending` which determine which column and which order we would like to sort by.

```
# Find the ten most northerly cities in Oregon  
cities.sort_values(by='lat', ascending=False).head(10)
```

# Sorting

We can sort `DataFrames` using the `sort_values` function. This function takes two arguments, `by` and `ascending` which determine which column and which order we would like to sort by.

```
# Find the ten most northerly cities in Oregon
cities.sort_values(by='lat', ascending=False).head(10)
```


	name	lat	lon	geometry
13	Astoria	46.19	-123.81	POINT (-123.81000 46.19000)
354	Warrenton	46.17	-123.92	POINT (-123.92000 46.17000)
159	Jeffers Gardens	46.15	-123.85	POINT (-123.85000 46.15000)
363	Westport	46.13	-123.37	POINT (-123.37000 46.13000)
60	Clatskanie	46.10	-123.21	POINT (-123.21000 46.10000)
269	Rainier	46.09	-122.95	POINT (-122.95000 46.09000)



# Geometric properties

The special thing about a **GeoDataFrame** is that it contains a **geometry** column. We can therefore apply spatial methods to these data. To demonstrate we will use our Oregon county shapefile.

```
# Read shapefile  
counties = gpd.read_file('data/orcntypoly.shp')
```



# Projections

**GeoDataFrames** have their own **CRS** which can be accessed using the **crs** method. The CRS tells **GeoPandas** where the coordinates of the geometries are located on the Earth's surface.

```
counties.crs
```

# Projections

**GeoDataFrames** have their own **CRS** which can be accessed using the **crs** method. The CRS tells **GeoPandas** where the coordinates of the geometries are located on the Earth's surface.

```
counties.crs
```

```
<Compound CRS: EPSG:5498>
Name: NAD83 + NAVD88 height
Axis Info [ellipsoidal|vertical]:
- Lat[north]: Geodetic latitude (degree)
- Lon[east]: Geodetic longitude (degree)
- H[up]: Gravity-related height (metre)
Area of Use:
- name: United States (USA) – CONUS and Alaska – onshore – Alabama; Alaska mainland; Arizona; .
- bounds: (-168.26, 24.41, -66.91, 71.4)
Datum: North American Datum 1983
- Ellipsoid: GRS 1980
- Prime Meridian: Greenwich
Sub CRS:
- NAD83
- NAVD88 height
```

# Projections

We can reproject a our data using the `to_crs` method.

```
counties_reproject = counties.to_crs('EPSG:32610')  
counties_reproject.crs
```

<Derived Projected CRS: EPSG:32610>

Name: WGS 84 / UTM zone 10N

Axis Info [cartesian]:

- E[east]: Easting (metre)
- N[north]: Northing (metre)

Area of Use:

- name: Between 126°W and 120°W, northern hemisphere between equator and 84°N, onshore and off
- bounds: (-126.0, 0.0, -120.0, 84.0)

Coordinate Operation:

- name: UTM zone 10N
- method: Transverse Mercator

Datum: World Geodetic System 1984 ensemble


- Ellipsoid: WGS 84
- Prime Meridian: Greenwich

# Geometric properties

Now our data has a **projected CRS**, we can calculate the area of each county with no warnings.

```
counties_reproject['area'] = counties_reproject['geometry'].area
counties_reproject.head()
```

	county	geometry	area
0	Josephine County	POLYGON ((481193.348 4727817.470, 481193.815 4...	4.246561e+09
1	Curry County	POLYGON ((433624.799 4737686.263, 433624.695 4...	5.162416e+09
2	Jackson County	POLYGON ((558466.270 4760676.349, 558468.456 4...	7.249726e+09
3	Coos County	POLYGON ((433624.799 4737686.263, 433230.247 4...	4.684074e+09
4	Klamath County	POLYGON ((634510.330 4830646.470, 634516.612 4...	1.588783e+10




# Geometric properties

There are other spatial methods we can apply to polygons such as the length of the outer edge (i.e. perimeter).

```
counties_reproject['perimeter'] = counties_reproject['geometry'].length
counties_reproject.head()
```

	county	geometry	area	perimeter
0	Josephine County	POLYGON ((481193.348 4727817.470, 481193.815 4...	4.246561e+09	331219.627981
1	Curry County	POLYGON ((433624.799 4737686.263, 433624.695 4...	5.162416e+09	438569.090594
2	Jackson County	POLYGON ((558466.270 4760676.349, 558468.456 4...	7.249726e+09	378774.110802
3	Coos County	POLYGON ((433624.799 4737686.263, 433230.247 4...	4.684074e+09	341177.961347
4	Klamath County	POLYGON ((634510.330 4830646.470, 634516.612 4...	1.588783e+10	619485.806580



# Geometric properties


Our `cities GeoDataFrame` also has geometric properties. We can access the latitude and longitude using the `x` and `y` methods.

```
cities['geometry'].x
```

```
0    -123.22  
1    -118.56  
2    -117.07  
3    -123.10  
4    -122.87
```

```
cities['geometry'].y
```

```
0     44.67  
1     45.77  
2     43.74  
3     44.63  
4     45.49
```



# Measure distance

We can measure the distance between two points, provided they have a projected CRS.

```
cities_reproject = cities.to_crs('EPSG:32610')
```



# Measure distance

We can measure the distance between two points, provided they have a projected CRS.

```
cities_reproject = cities.to_crs('EPSG:32610')
```

```
eugene = cities_reproject[cities_reproject['name'] == 'Eugene'].reset_index()  
bend = cities_reproject[cities_reproject['name'] == 'Bend'].reset_index()
```

# Measure distance

We can measure the distance between two points, provided they have a projected CRS.

```
cities_reproject = cities.to_crs('EPSG:32610')
```

```
eugene = cities_reproject[cities_reproject['name'] == 'Eugene'].reset_index()  
bend = cities_reproject[cities_reproject['name'] == 'Bend'].reset_index()
```


```
eugene.distance(bend).values[0] / 1000
```

```
144.97607871968486
```

# Plot

If our data are in the same projection system, we can plot them together.

```
ax = counties_reproject.plot(figsize=(10, 10), alpha=0.5, edgecolor='k')  
cities_reproject.plot(ax=ax, color='black', markersize=5)
```



# Spatial joins

One of the most useful things about `GeoPandas` is that it contains functions to perform **spatial joins** to combine two `GeoDataFrames` based on the **spatial relationships** between their geometries.

The order of the two `GeoDataFrames` is quite important here, as well as the `how` argument. A **left** outer join implies that we are interested in **retaining the geometries** of the `GeoDataFrame` on the left, i.e. the **point** locations of the cities. We then retain attributes of the **right** `GeoDataFrame` if they intersect and drop them if they don't.

# Spatial joins

The following would provide the county attributes for all our cities based on a spatial intersection.

```
cities_reproject.sjoin(counties_reproject, how="left").head()
```

	name	lat	lon	geometry	index_right	county	ar
0	Adair Village city	44.67	-123.22	POINT (482561.392 4946316.184)	12	Benton County	1.755207e+
1	Adams	45.77	-118.56	POINT (845212.127 5078087.252)	33	Umatilla County	8.385719e+
2	Adrian	43.74	-117.07	POINT (977541.425 4860113.062)	10	Malheur County	2.581882e+
3	Albany	44.63	-123.10	POINT (492067.910 4941854.290)	13	Linn County	5.969370e+
4	Aloha	45.49	-122.87	POINT (510158.282 5037393.753)	27	Washington County	1.880526e+

# Which county contains the most cities/towns?

We would do this would be to use the `groupby` function.

```
join = cities_reproject.sjoin(counties_reproject, how="left")
```

# Which county contains the most cities/towns?

We would do this would be to use the `groupby` function.

The first argument `groupby` accepts is the column we want group our data into (`county` in our case). Next, it takes a column (or list of columns) to summarize. Finally, this function does nothing until we specify **how** we want to group our data (`count`).

It's actually nice to **reset the index** after using `groupby` so that we end up with a DataFrame (rather than a Series).

```
grouped = join.groupby('county')['name'].count().reset_index()  
grouped.nlargest(n=10, columns='name')
```

	county	name
23	Marion County	25
33	Washington County	24
2	Clackamas County	23
9	Douglas County	23
21	Linn County	23
29	Umatilla County	19
28	Tillamook County	18
14	Jackson County	17
8	Deschutes County	14
19	Lane County	12

# Next time: Network data analysis



Email: [jryan4@uoregon.edu](mailto:jryan4@uoregon.edu)

Office: 163A Condon Hall

Office hours: Monday 15:00-16:00 and Tuesday 14:00-15:00